

# Feature Engineering (Vectorization)

---

PREPARED BY: AHMAD ALAA ALDINE

PRESENTED BY: AHMAD ALAA ALDINE





What We See

08 02 22 97 38 15 00 40 00 75 04 05 07 78 52 12 50 77 91 08  
49 49 99 40 17 81 18 57 60 87 17 40 98 43 69 48 04 56 62 00  
81 49 31 73 55 79 14 29 93 71 40 67 53 88 30 03 49 13 36 65  
52 70 95 23 04 60 11 42 69 24 68 56 01 32 56 71 37 02 36 91  
22 31 16 71 51 67 63 89 41 92 36 54 22 40 40 28 66 33 13 80  
24 47 32 60 99 03 45 02 44 75 33 53 78 36 84 20 35 17 12 50  
32 98 81 28 64 23 67 10 26 38 40 67 59 54 70 66 18 38 64 70  
67 26 20 68 02 62 12 20 95 63 94 39 63 08 40 91 66 49 94 21  
24 55 58 05 66 73 99 26 97 17 78 78 96 83 14 88 34 89 63 72  
21 36 23 09 75 00 76 44 20 45 35 14 00 61 33 97 34 31 33 95  
78 17 53 28 22 75 31 67 15 94 03 80 04 62 16 14 09 53 56 92  
16 39 05 42 96 35 31 47 55 58 88 24 00 17 54 24 36 29 85 57  
86 56 00 48 35 71 89 07 05 44 44 37 44 60 21 58 51 54 17 58  
19 80 81 68 05 94 47 69 28 73 92 13 86 52 17 77 04 89 55 40  
04 52 08 83 97 35 99 16 07 97 57 32 16 26 26 79 33 27 98 66  
88 36 68 87 57 62 20 72 03 46 33 67 46 55 12 32 63 93 53 69  
04 42 16 73 38 25 39 11 24 94 72 18 08 46 29 32 40 62 76 36  
20 69 36 41 72 30 23 88 34 62 99 69 82 67 59 85 74 04 36 16  
20 73 35 29 78 31 90 01 74 31 49 71 48 86 81 16 23 57 05 54  
01 70 54 71 83 51 54 69 16 92 33 48 61 43 52 01 89 19 67 48

What Computers See

# Definition

---

Feature engineering in Natural Language Processing (NLP) involves **transforming** raw text data into a **numerical** format that can be effectively utilized by **machine learning** algorithms.

- I. Document to vector (Document Embedding)
  - a. Bag of Words – Count Vectorization
  - b. Bag of Words – TF-IDF
- II. Word to vector (Word Embedding)
  - a. Distributional Word Representation
  - b. Word2vec

# Bag of Words

---

# Bag of Words (1/3)

---

The bag-of-words model is a way of representing text data when modeling text with machine learning algorithms.

The bag-of-words model is simple to understand and implement and has seen great success in problems such as language modeling, document classification, sentiment analysis, and others.

A bag-of-words is a representation of text that describes the occurrence of words within a document. It involves two things:

1. A vocabulary of known words
2. A measure of the presence of known words

# Bag of Words (2/3)

---

## BoW Steps:

1. Collect Dataset
  - Set of documents
2. Desing the vocabulary
  - Set of unique words in the dataset (ignoring stopwords and punctuations)
3. Create document vectors
  - Turn each document into a vector that we can use in machine learning

# Bag of Words (3/3)

---

## Limitations:

### 1. Semantic meaning

- It does not consider the semantic meaning of a word
- It ignores the context in which the word is used

### 2. Sparse Vectorization

- Large feature space with many zeros
- Higher computational time

# Bag of Words – Count Vectorization

**Example:**

**Dataset**

D1	The dog barked in the park on another dog.
D2	The owner of the dog put him on the leash since he barked.
D3	My dog is barking and chasing its tail.

**Vocabulary**

{'chasing', 'leash', 'barked', 'tail', 'barking', 'park', 'owner', 'dog'}

**Count  
Vectorization**

	chasing	leash	barked	tail	barking	park	owner	dog
D1	0	0	1	0	0	1	0	2
D2	0	1	1	0	0	0	1	1
D3	1	0	0	1	1	0	0	1



# Bag of Words – TF-IDF (1/3)

---

TF-IDF stands for **Term Frequency — Inverse Document Frequency**

If a particular word appears multiple times in a document, then it might have higher importance than the other words that appear fewer times (TF).

At the same time, if a particular word appears many times in a document, but it is also present many times in some other documents, then maybe that word is frequent, so we cannot assign much importance to it. (IDF).

# TF-IDF (2/3)

---

## Formula:

$$TF - IDF = TF * IDF$$

$$TF = \frac{\text{Frequency of the word in the document}}{\text{Total number of words in the document}}$$

$$IDF = \log \left( \frac{\text{Total number of documents}}{\text{Number of documents containing the word}} \right)$$

# TF-IDF(3/3)

Example:

Count Vectorization

	chasing	leash	barked	tail	barking	park	owner	dog
D1	0	0	1	0	0	1	0	2
D2	0	1	1	0	0	0	1	1
D3	1	0	0	1	1	0	0	1

TF

	chasing	leash	barked	tail	barking	park	owner	dog
D1	0	0	$\frac{1}{4}=0.25$	0	0	$\frac{1}{4}=0.25$	0	$\frac{2}{4}=0.5$
D2	0	$\frac{1}{4}=0.25$	$\frac{1}{4}=0.25$	0	0	0	$\frac{1}{4}=0.25$	$\frac{1}{4}=0.25$
D3	$\frac{1}{4}=0.25$	0	0	$\frac{1}{4}=0.25$	$\frac{1}{4}=0.25$	0	0	$\frac{1}{4}=0.25$

IDF

	chasing	leash	barked	tail	barking	park	owner	dog
IDF	$\text{Log}(3/1)=0.477$	$\text{Log}(3/1)=0.477$	$\text{Log}(3/2)=0.176$	$\text{Log}(3/1)=0.477$	$\text{Log}(3/1)=0.477$	$\text{Log}(3/1)=0.477$	$\text{Log}(3/1)=0.477$	$\text{Log}(3/3)=0$

TF-IDF

	chasing	leash	barked	tail	barking	park	owner	dog
D1	0	0	$0.25*0.176$	0	0	$0.25*0.477$	0	0
D2	0	$0.25*0.477$	$0.25*0.176$	0	0	0	$0.25*0.477$	0
D3	$0.25*0.477$	0	0	$0.25*0.477$	$0.25*0.477$	0	0	0

# Distributional Word Representation

---

# Distributional Word Representation (1/2)

---

Represent words as **vectors** based on the distribution of their **context** words.

Words sharing the **same context** tend to have a similar **meaning**.

Steps:

1. Define what are the contexts of a word in the dataset
  - Window (surrounding words)
  - All words in a sentence
2. Count how many times each vocabulary word occurs in these contexts
3. Build vectors

# Distributional Word Representation (2/3)

**Example:**

**Dataset**

D1	The dog barked in the park on another dog.
D2	The owner of the dog put him on the leash since he barked.
D3	My dog is barking and chasing its tail.

**Vocabulary**

{'chasing', 'leash', 'barked', 'tail', 'barking', 'park', 'owner', 'dog'}

**Vectors  
(Window = 5)**

	chasing	leash	barked	tail	barking	park	owner	dog
chasing	0	0	0	1	1	0	0	0
leash	0	0	0	0	0	0	0	0
barked	0	0	0	0	0	0	0	1
tail	1	0	0	0	0	0	0	0
barking	1	0	0	0	0	0	0	1
park	0	0	0	0	0	0	0	0
owner	0	0	0	0	0	0	0	0
dog	0	0	0	0	1	0	0	0

# Distributional Word Representation (3/3)

**Example:**

**Dataset**

D1	The dog barked in the park on another dog.
D2	The owner of the dog put him on the leash since he barked.
D3	My dog is barking and chasing its tail.

**Vocabulary**

{'chasing', 'leash', 'barked', 'tail', 'barking', 'park', 'owner', 'dog'}

**Vectors  
(all sentence)**

	chasing	leash	barked	tail	barking	park	owner	dog
chasing	0	0	0	1	1	0	0	1
leash	0	0	1	0	0	0	1	1
barked	0	1	0	0	0	1	1	3
tail	1	0	0	0	1	0	0	1
barking	1	0	0	1	0	0	0	1
park	0	0	1	0	0	1	0	1
owner	0	1	1	0	0	0	0	1
dog	1	1	2	1	1	0	0	0

**Large dimension matrix with many zeros → Sparsity Problem**

# Word2Vec

---



# Word2Vec

---

This approach was released back in 2013 by **Google** researchers, and it took the NLP industry by storm.

It uses the power of a simple **Neural Network** to generate word embeddings.

There are mainly two ways to implement Word2Vec:

1. Skip-Gram
2. CBOW

# Word2Vec – Skip-Gram (1/4)


In Skip-Gram method, we provide a word to our Neural Network and ask it to predict the context.

Technically, it predicts the probabilities of a word being a context word for the given target word.

However, the interesting part is, we don't use this trained Neural Network.

Instead, the goal is just to learn the weights of the hidden layer while predicting the surrounding words correctly.

These weights are the **word embeddings**.



# Word2Vec – Skip-Gram (2/4)

## Input Layer:

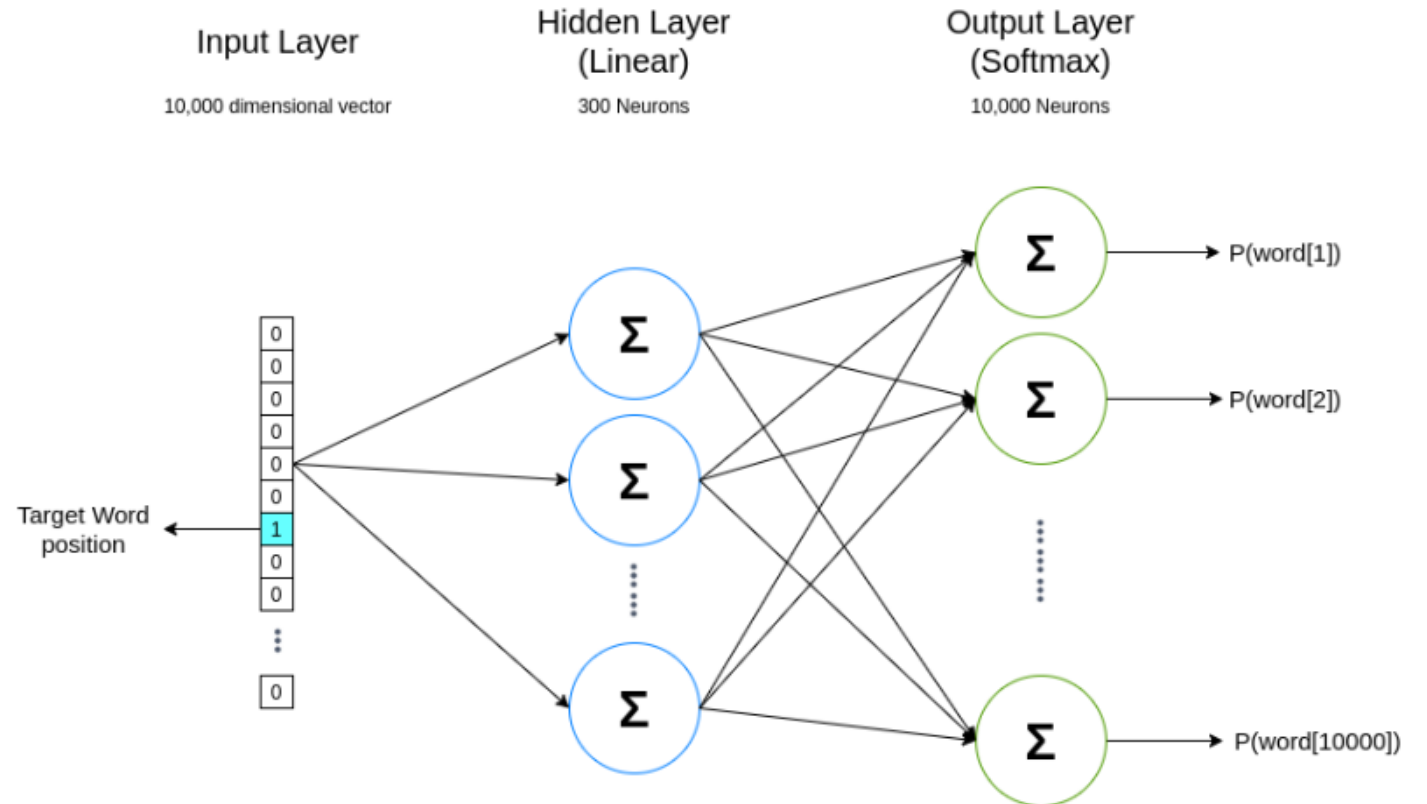
- Dataset vocabulary words

## Hidden Layer:

- Hyperparameter
- tuned to obtain the best results

## Output Layer:

- Dataset vocabulary words



# Word2Vec – Skip-Gram (3/4)

Input:

- One-Hot Encoding for a target word
- 1 in the position corresponding to the target word and 0 everywhere

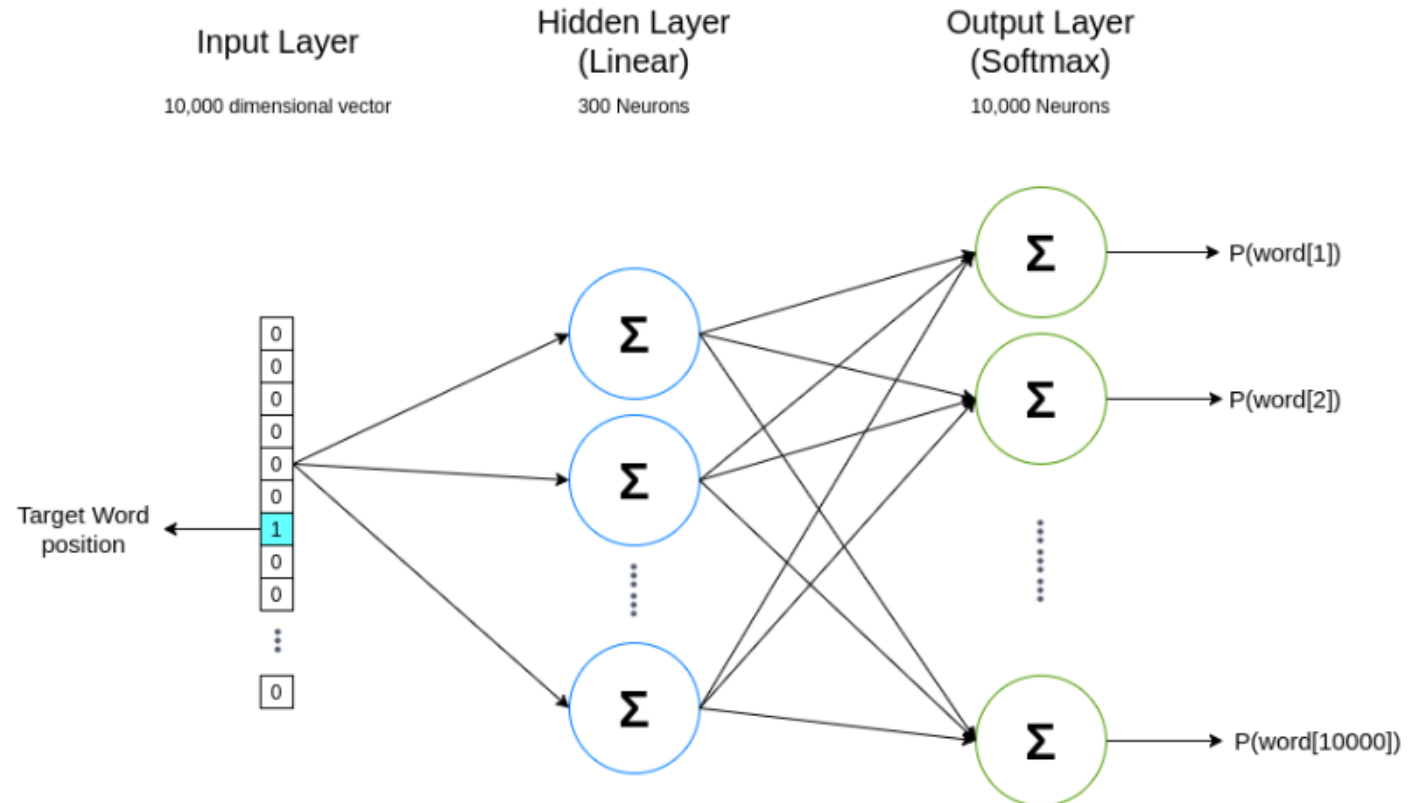
Hidden Layer:

- Linear (No activation function)

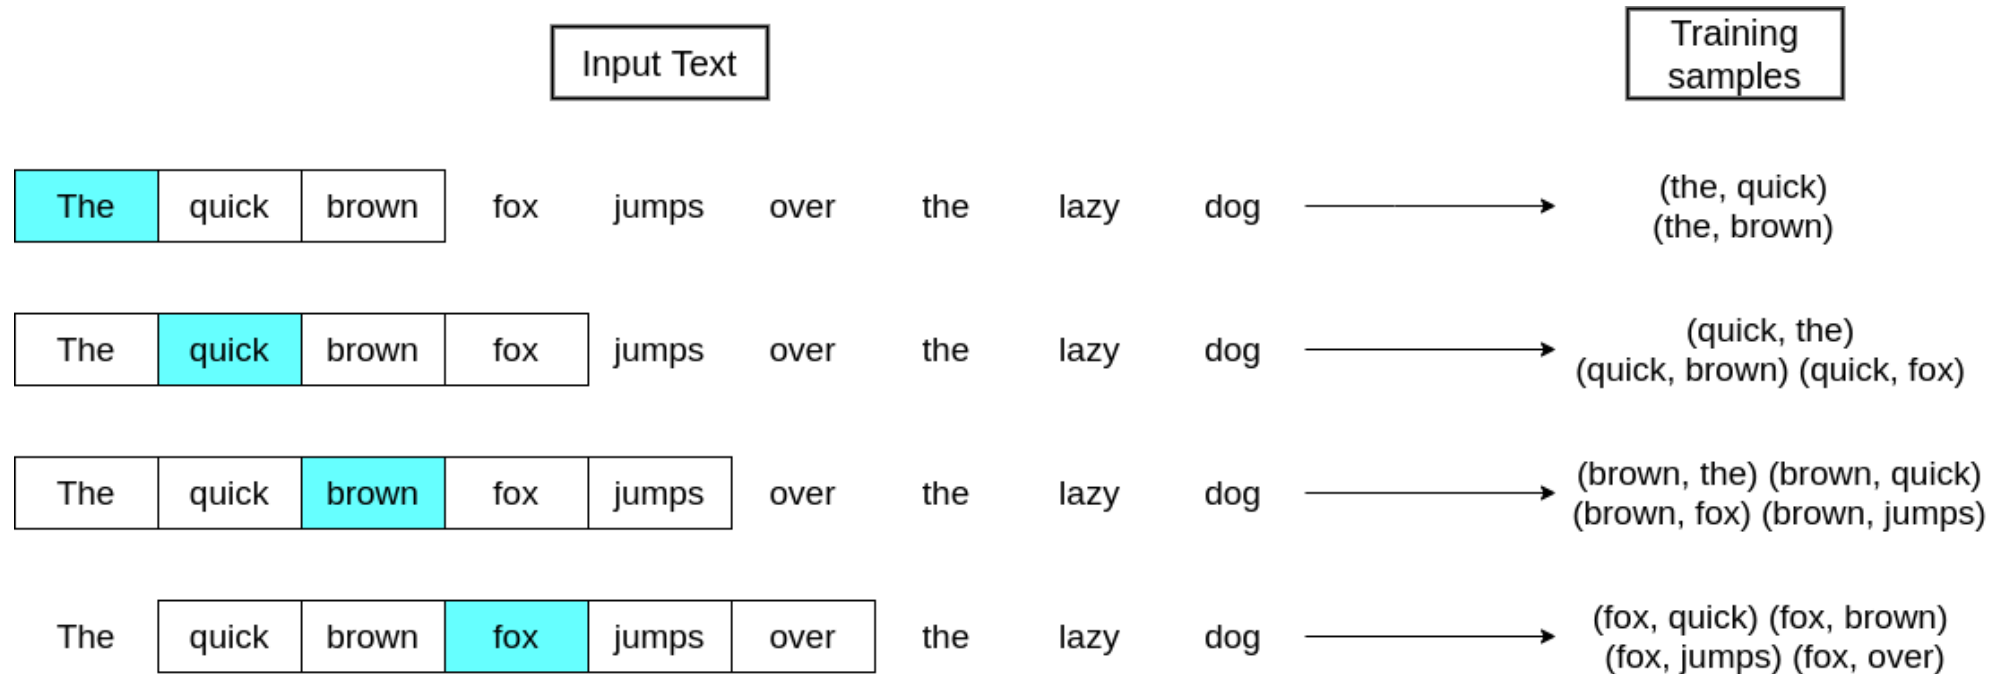
Output:

- The probability of each word being the context word for our input target word

- Softmax:  $S(y_i) = \frac{e^{y_i}}{\sum_{j=1}^N e^{y_j}}$



# Word2Vec – Skip-Gram (4/4)




Once the model gets trained, the final word embedding for each word will be given by the following calculation: **1×10000 input vector \* 10000×300 weights matrix = 1×300 vector**

# Word2Vec – CBOW (1/2)

CBOW stands for Continuous Bag of Words.

Instead of predicting the context words, we input them into the model and ask the network to predict the current word.

CBOW is the mirror image of the skip-gram approach.



# Word2Vec – CBOW (2/2)


The dimension of our hidden layer and output layer stays the same as the skip-gram model.

The input is  $C$  context words in the form of a one-hot encoded vector of size  $1 \times V$  ( $V$  = size of vocabulary).

$C$  vectors will be multiplied by the Weights of our hidden layer of shape  $V \times N$  ( $N$  = number of neurons in the hidden layer).

This will result in  $C$ ,  $1 \times N$  vectors, and all of these  $C$  vectors will be averaged element-wise to obtain our final activation for the hidden layer, which then will be fed into our output softmax layer.

The learned weight between the hidden and output layer makes up the word embedding representation.



# Word2Vec – Skip-Gram vs CBOW


---

When to use the skip-gram model and when to use CBOW?


- Skip-gram works well with small datasets and can better represent rare words.
- CBOW is found to train faster than skip-gram and can better represent frequent words.
- So the choice of skip-gram VS. CBOW depends on the kind of problem that we're trying to solve.




# Word2Vec – Results



Male-Female




Verb tense



Country-Capital

# Word2Vec – Results



# Feature Engineering in Python

---

## Count Vectorization using Sklearn

```
from sklearn.feature_extraction.text import CountVectorizer

sents = ['The dog barked in the park on another dog',
         'The owner of the dog put him on the leash since he barked.',
         'My dog is barking and chasing its tail.']

cv = CountVectorizer()

X = cv.fit_transform(sents)
X = X.toarray()

print(sorted(cv.vocabulary_.keys()))
print(X)
```

```
['and', 'another', 'barked', 'barking', 'chasing', 'dog', 'he', 'him', 'in', 'is', 'its', 'leash', 'my', 'of', 'on', 'owner',
'park', 'put', 'since', 'tail', 'the']
[[0 1 1 0 0 2 0 0 1 0 0 0 0 0 1 0 1 0 0 0 2]
 [0 0 1 0 0 1 1 1 0 0 0 1 0 1 1 1 0 1 1 0 3]
 [1 0 0 1 1 1 0 0 0 1 1 0 1 0 0 0 0 0 0 1 0]]
```

## Count Vectorization without Stopwords using Sklearn

```
from sklearn.feature_extraction.text import CountVectorizer

sents = ['The dog barked in the park on another dog',
         'The owner of the dog put him on the leash since he barked.',
         'My dog is barking and chasing its tail.']

cv = CountVectorizer(stop_words='english')

X = cv.fit_transform(sents)
X = X.toarray()

print(sorted(cv.vocabulary_.keys()))
print(X)

['barked', 'barking', 'chasing', 'dog', 'leash', 'owner', 'park', 'tail']
[[1 0 0 2 0 0 1 0]
 [1 0 0 1 1 1 0 0]
 [0 1 1 1 0 0 0 1]]
```

## N-Gram Count Vectorization without Stopwords using Sklearn ¶

```
from sklearn.feature_extraction.text import CountVectorizer

sents = ['The dog barked in the park on another dog',
         'The owner of the dog put him on the leash since he barked.',
         'My dog is barking and chasing its tail.']

cv = CountVectorizer(ngram_range=(1,2), stop_words='english')

X = cv.fit_transform(sents)
X = X.toarray()

print(sorted(cv.vocabulary_.keys()))
print(X)
```

```
['barked', 'barked park', 'barking', 'barking chasing', 'chasing', 'chasing tail', 'dog', 'dog barked', 'dog barking', 'dog leash', 'leash', 'leash barked', 'owner', 'owner dog', 'park', 'park dog', 'tail']
[[1 1 0 0 0 0 2 1 0 0 0 0 0 0 1 1 0]
 [1 0 0 0 0 0 1 0 0 1 1 1 1 1 0 0 0]
 [0 0 1 1 1 1 1 0 1 0 0 0 0 0 0 0 1]]
```

## TF-IDF Vectorization using Sklearn

```
from sklearn.feature_extraction.text import TfidfVectorizer

sents = ['The dog barked in the park on another dog',
         'The owner of the dog put him on the leash since he barked.',
         'My dog is barking and chasing its tail.']

cv = TfidfVectorizer(stop_words='english')

X = cv.fit_transform(sents)

print(X.toarray())

import pandas as pd
df = pd.DataFrame(X[0].T.todense(), index=cv.get_feature_names(), columns=["TF-IDF"])
df = df.sort_values('TF-IDF', ascending=False)

print()
print(df)
```

```
[[0.44102652 0.          0.68499287 0.          0.
  0.57989687 0.          ]
 [0.44451431 0.          0.34520502 0.5844829  0.5844829
  0.          0.          ]
 [0.          0.54645401 0.54645401 0.32274454 0.          0.
  0.          0.54645401]]
```

	TF-IDF
dog	0.684993
park	0.579897
barked	0.441027
barking	0.000000
chasing	0.000000
leash	0.000000
owner	0.000000
tail	0.000000

## Word2Vec: using Pre-trained model in Gensim

```
from gensim import models
w2v = models.KeyedVectors.load_word2vec_format(
    './GoogleNews-vectors-negative300.bin', binary=True)

vect = w2v['healthy']
w2v.most_similar('happy')
```

## Word2Vec: Train our model in Gensim

```
from gensim import models

sents = ['The dog barked in the park on another dog',
         'The owner of the dog put him on the leash since he barked.',
         'My dog is barking and chasing its tail.']

# Word2vec requires the training dataset in form of a list of lists
# of tokenized sentences, so we'll preprocess and convert sents to

sents = [sent.split() for sent in sents]
custom_model = models.Word2Vec(sents, min_count=1, vector_size=300, workers=4)

vect = custom_model.wv['dog']
print(vect[:5])
result = custom_model.wv.most_similar('barked')
print(result)
```

```
[-1.7874241e-04  7.8810059e-05  1.7011166e-03  3.0030911e-03
 -3.1009833e-03]
[('and', 0.11018942296504974), ('leash', 0.09745844453573227), ('him', 0.08689301460981369), ('barking', 0.07596393674612045),
 ('he', 0.07124919444322586), ('put', 0.06887509673833847), ('owner', 0.053101178258657455), ('its', 0.03030511364340782), ('do
g', 0.026307500898838043), ('is', 0.023843316361308098)]
```